

Instructor: Prof. Onur Mutlu

TAs: Juan Gomez Luna, Mohammad Sadrosadati, Mohammed Alser, Ataberk Olgun, Giray Yaglikci, Can Firtina, Geraldo F. de Oliveira Junior, Rahul Bera, Konstantinos Kanellopoulos, Nika Mansouri Ghiasi, Nisa Bostancı, Rakesh Nadig, Joel Lindegger, İsmail Emir Yüksel, Haocong Luo, Yahya Can Tuğrul, Julien Eudine

Released: Monday, April 24, 2023

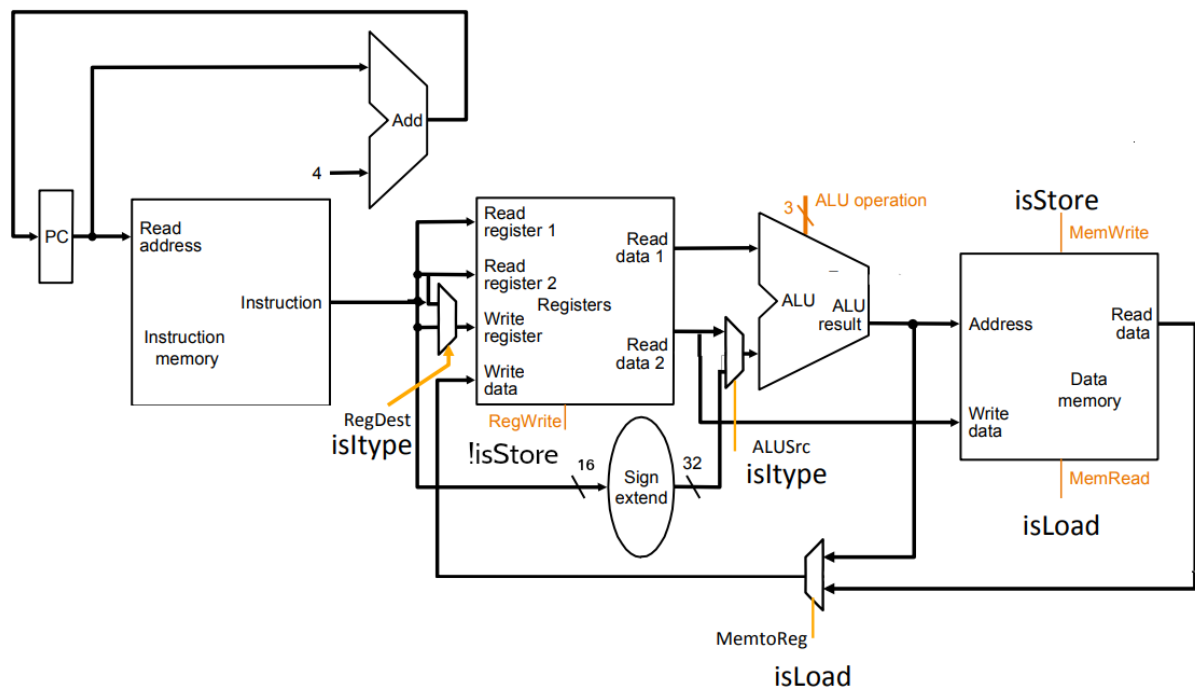
1 Single-Cycle Processor Datapath

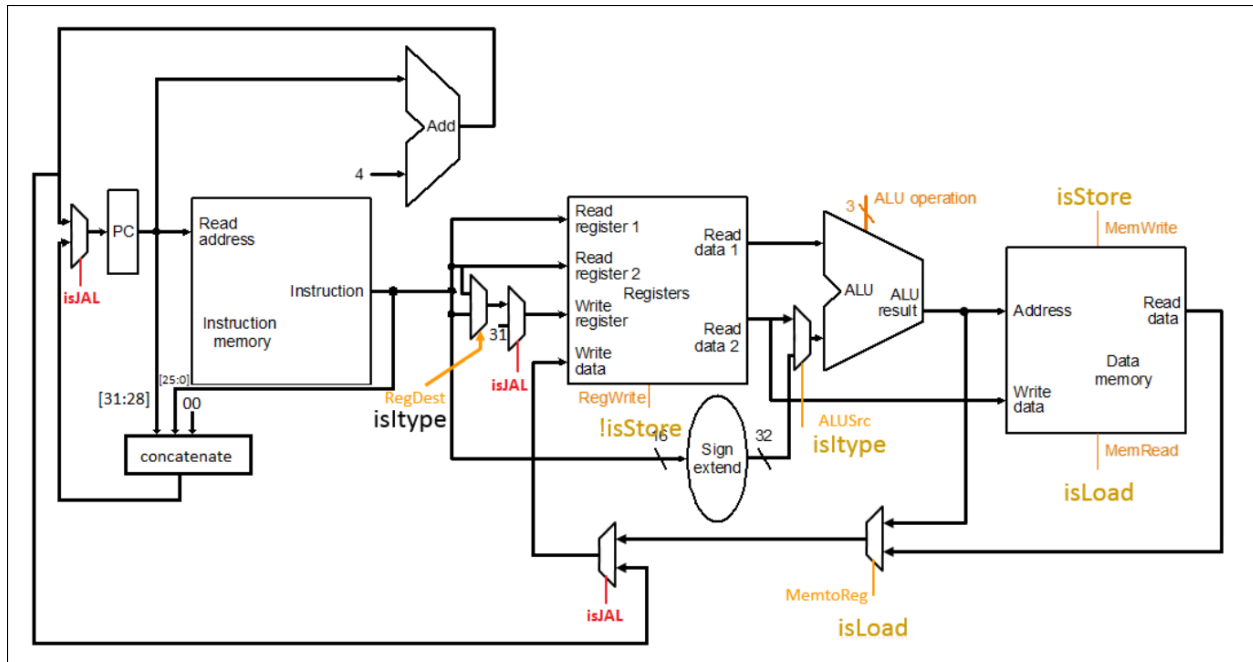
In this problem, you will modify the single-cycle datapath we built up in Lecture 11 to support the **JAL** instruction. The datapath that we will start with is provided below. Your job is to implement the necessary data and control signals to support the **JAL** instruction, which we define to have the following semantics:

JAL : $R31 \leftarrow PC + 4$

$$\text{PC} \leftarrow \text{PC}_{31 \dots 28} \parallel \text{Immediate} \parallel 0^2$$

Add to the datapath on the next page the necessary data and control signals to implement the **JAL** instruction. Draw and label all components and wires very clearly (give control signals meaningful names; if selecting a subset of bits from many, specify exactly which bits are selected; and so on).





2 REP MOVSB

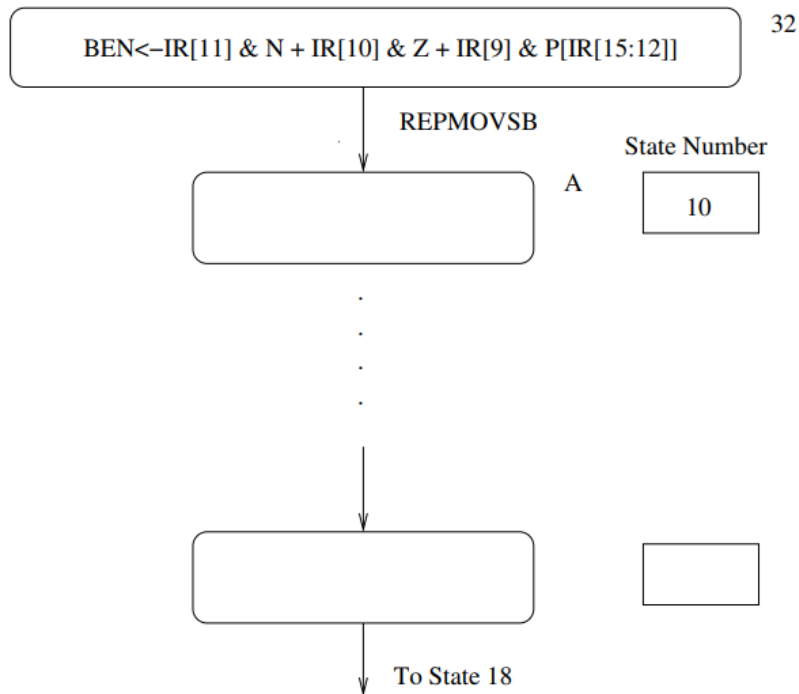
Let's say you are the lead architect of the next flagship processor at Advanced Number Devices (AND). You have decided that you want to use the LC-3b ISA for your next product, but your customers want a smaller semantic gap and marketing is on your case about it. So, you have decided to implement your favorite x86 instruction, REP MOVSB, in LC-3b.

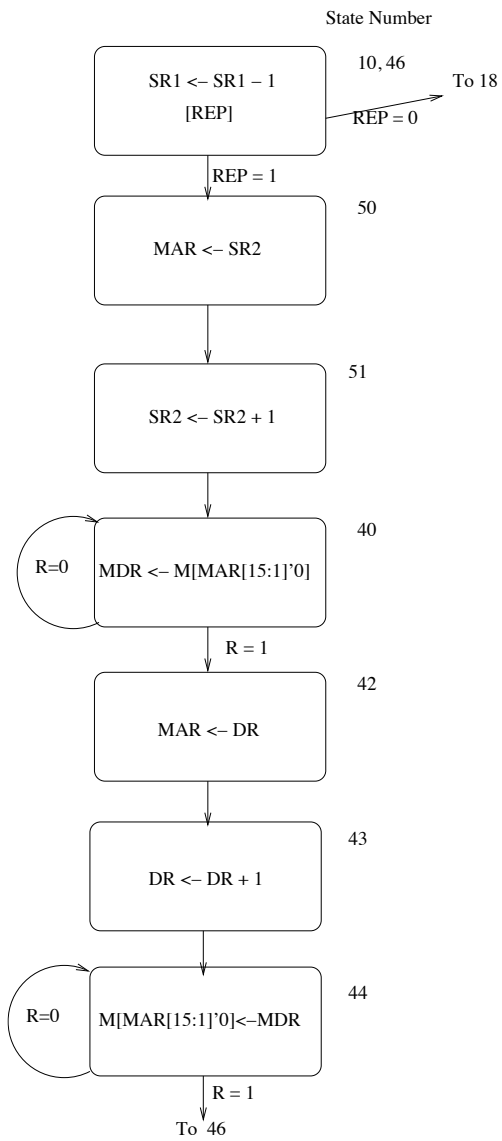
Specifically, you want to implement the following definition for REP MOVSB (in LC-3b parlance): REP-MOVSB SR1, SR2, DR which is encoded in LC-3b machine code as:

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
1010				DR			SR1			0	0	0	SR2		

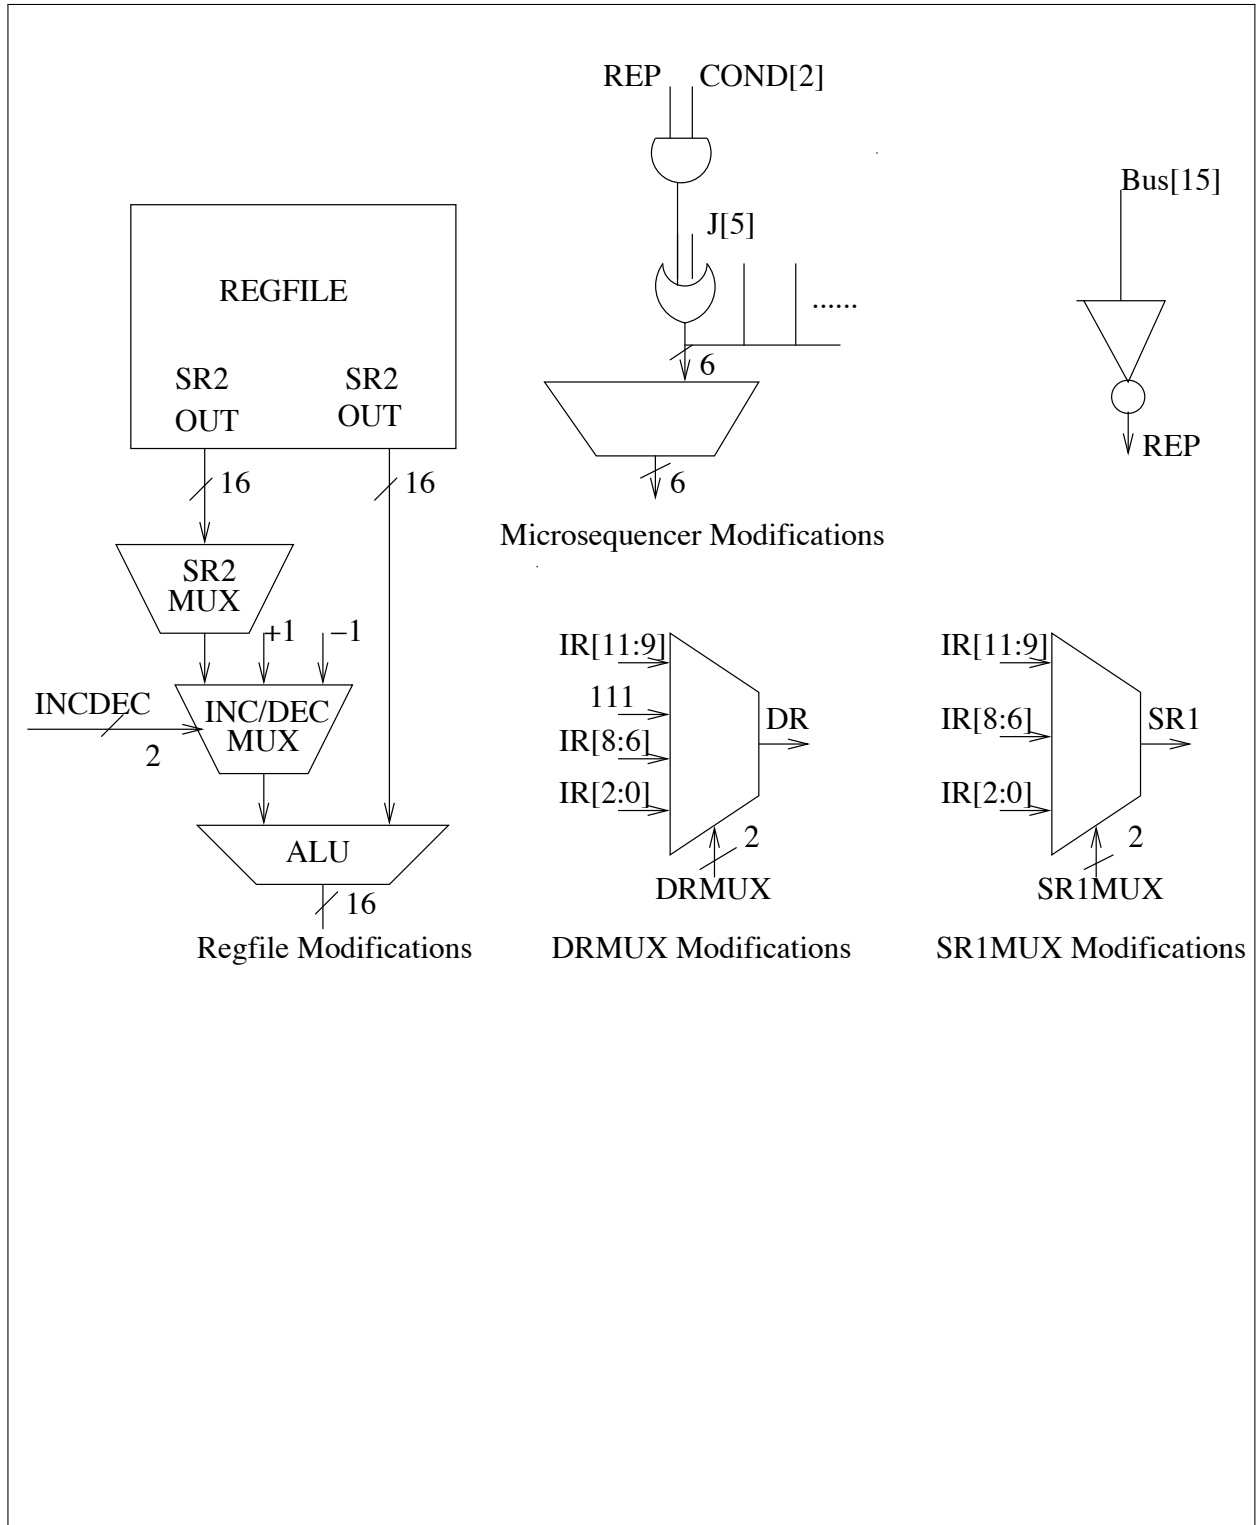
REPMOVSB uses three registers: SR1 (count), SR2 (source), and DR (destination). It moves a byte from memory at address SR2 to memory at address DR, and then increments both pointers by one. This is repeated SR1 times. Thus, the instruction copies SR1 bytes from address SR2 to address DR. Assume that the value in SR1 is greater than or equal to zero.

1. Complete the state diagram shown below, using the notation of the LC-3b state diagram. Describe inside each bubble what happens in each state and assign each state an appropriate state number. Add additional states not present in the original LC-3b design as you see fit.





2. Add to the LC-3b datapath any additional structures and any additional control signals needed to implement REPMOVSb. Clearly label your additional control signals with descriptive names. Describe what value each control signal would take to control the datapath in a particular way.



3. Describe any changes you need to make to the LC-3b microsequencer. Add any additional logic and control signals you need. Clearly describe the purpose and function of each signal and the values it would take to control the microsequencer in a particular way.

Additional control signals

- INCDEC/2: PASSSR2, +1, -1
- DRMUX/2:
 - IR[11:9] ;destination IR[11:9]
 - R7 ;destination R7
 - IR[8:6] ;destination IR[8:6]
 - IR[2:0] ;destination IR[2:0]
- SR1MUX/2:
 - IR[11:9] ;source IR[11:9]
 - IR[8:6] ;source IR[8:6]
 - IR[2:0] ;source IR[2:0]
- COND/3:
 - COND0: Unconditional
 - COND1: Memory Ready
 - COND2: Branch
 - COND3: Addressing Mode
 - COND4: Repeat

3 Pipelining

Consider two pipelined machines implementing MIPS ISA, Machine I and Machine II:

Both machines have the following *five pipeline stages*, very similarly to the basic 5-stage pipelined MIPS processor we discussed in lectures, and *one ALU*:

1. Fetch (one clock cycle)
2. Decode (one clock cycle)
3. Execute (one clock cycle)
4. Memory (one clock cycle)
5. Write-back (one clock cycle).

Machine I does *not* implement interlocking in hardware. It assumes all instructions are independent and relies on the compiler to order instructions such that there is sufficient distance between dependent instructions. The compiler either moves other independent instructions between two dependent instructions, if it can find such instructions, or otherwise, inserts **nops**. Assume internal register file forwarding (an instruction writes into a register in the first half of a cycle and another instruction can correctly access the same register in the next half of the cycle). Assume that the processor predicts all branches as always-taken.

Machine II implements data forwarding in hardware. On detection of a flow dependence, it forwards an operand from the memory stage or from the write-back stage to the execute stage. The load instruction (**lw**) can *only* be forwarded from the write-back stage because data becomes available in the memory stage but not in the execute stage like for the other instructions. Assume internal register file forwarding (an instruction writes into a register in the first half of a cycle and another instruction can access the same register in the next half of the cycle). The compiler does *not* reorder instructions. Assume that the processor predicts all branches as always-taken.

Consider the following code segment:

```
Copy: lw    $2, 100($5)
      sw    $2, 200($6)
      addi  $1, $1, 1
      bne   $1, $25, Copy
```

Initially, $\$5 = 0$, $\$6 = 0$, $\$1 = 0$, and $\$25 = 25$.

- (a) When the given code segment is executed on Machine I, the compiler has to reorder instructions and insert `nops` if needed. Write the resulting code that has minimal modifications from the original.

```
Copy:  lw $2, 100($5)
        addi $1, $1, 1
        nop
        sw $2, 200($6)
        bne $1, $25, Copy
```

- (b) When the given code segment is executed on Machine II, dependencies between instructions are resolved in hardware. Explain when data is forwarded and which instructions are stalled and when they are stalled.

In every iteration, data are forwarded for `sw` and for `bne`. The instruction `sw` is dependent on `lw`, so it is stalled one cycle in every iteration

- (c) Calculate the *machine code size* of the code segments executed on Machine I (part (a)) and Machine II (part (b)).

Machine I - 20 bytes (because of the additional `nop`)
Machine II - 16 bytes

- (d) Calculate the number of cycles it takes to execute the code segment on Machine I and Machine II.

Machine I: The compiler reorders instructions and places one `nop`. This is the execution timeline of the first iteration:

1	2	3	4	5	6	7	8	9
F	D	E	M	W				
	F	D	E	M	W			
		N	N	N	N	N		
			F	D	E	M	W	
				F	D	E	M	W

9 cycles for one iteration. As there are 5 instructions in each iteration and 25 iterations, the total number of cycles is 129 cycles.

Machine II: The machine stalls `sw` one cycle in the decode stage. This is the execution timeline of the first iteration:

1	2	3	4	5	6	7	8	9
F	D	E	M	W				
	F	D	D	E	M	W		
		F	F	D	E	M	W	
			F	D	E	M	W	

9 cycles for one iteration. As there are 4 instructions in each iteration and 25 iterations, and one stall cycle in each iteration, the total number of cycles is 129 cycles.

- (e) Which machine is faster for this code segment? Explain.

For this code segment, both machines take the same number of cycles. We cannot say which one is faster, since we do not know the clock frequency.

4 Pipeline - Reverse Engineering

The following piece of code runs on a pipelined microprocessor as shown in the table (F: Fetch, D: Decode, E: Execute, M: Memory, W: Write back). Instructions are in the form “Instruction Destination,Source1,Source2.” For example, “ADD A, B, C” means $A \leftarrow B + C$.

	Cycles	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18
0	MUL R5, R6, R7	F	D	E1	E2	E3	E4	M	W										
1	ADD R4, R6, R7		F	D	E1	E2	E3	-	M	W									
2	ADD R5, R5, R6			F	D	-	-	E1	E2	E3	M	W							
3	MUL R4, R7, R7				F	-	-	D	E1	E2	E3	E4	M	W					
4	ADD R6, R7, R5							F	D	-	E1	E2	E3	M	W				
5	ADD R3, R0, R6								F	-	D	-	-	E1	E2	E3	M	W	
6	ADD R7, R1, R4										F	-	-	D	E1	E2	E3	M	W

Use this information to reverse engineer the architecture of this microprocessor to answer the following questions. Answer the questions as precise as possible with the provided information. If the provided information is not sufficient to answer a question, answer “Unknown” and explain your reasoning clearly.

- (a) How many cycles does it take for an adder and for a multiplier to calculate a result?

3 cycles for adder (E1, E2, E3) and 4 cycles for multiplier (E1, E2, E3, E4).

- (b) What is the minimum number of register file read/write ports that this architecture implements? Explain.

The register file has two read ports and one write port.
Decode and Writeback stages can be performed simultaneously as seen at cycle 8. Decode reads from two registers, Writeback writes to one register.

- (c) Can we reduce the execution time of this code by enabling more read/write ports in the register file? Explain.

It is not possible to reduce stall cycles of the given code by only enabling more register file ports, as the pipeline would be stalled due to other limited resources.

- (d) Does this architecture implement any data forwarding? If so, how is data forwarding done between pipeline stages? Explain.

There is data forwarding from the M stage to E1, as we observe that the instruction 2 starts using R5 at the clk cycle 7, which is one clk cycle after the instruction 0 finishes calculating its result in the execution unit.

Similarly, as another proof of this data forwarding, we observe that the instruction 4 starts using R5 at the clk cycle 10, which is one clk cycle after the instruction 2 finishes calculating its result in the execution unit.

Any other data forwarding is *unknown* with the given information.

- (e) Is it possible to run this code faster by adding more data forwarding paths? If it is, how? Explain.

Not possible.

All instructions that stall due to data dependency are already using the best possible data forwarding. There is no stall cycles that can be eliminated by enabling another form of data forwarding.

- (f) Is there internal forwarding in the register file? If there is not, how would the execution time of the same program change by enabling internal forwarding in the register file? Explain.

The register file already implements internal forwarding, as instruction 6 can finish the decode stage by fetching the value of R4 from the register file in the same cycle that R4 is written (cycle 13).

- (g) Optimize the assembly code in order to reduce the number of stall cycles. You are allowed to *reorder*, *add*, or *remove* ADD and MUL instructions. You are expected to achieve the minimum possible execution time. Make sure that the register values that the optimized code generates at the end of its execution are identical to the register values that the original code generates at the end of its execution. Justify each individual change you make. Show the execution timeline of each instruction and what stage it is in the table below. (*Notice that the table below consists of two parts: the first ten cycles at the top, and the next ten cycles at the bottom.*)

- Instruction 1 is useless due to write-after-write, remove it.
- Instruction 3 stalls for decode logic, move it up.
- Instruction 6 does not have read-after-write dependency and can be executed before instr. 5. However, it cannot execute before instruction 4 as it would change the value of R7.

New total execution time is 17 cycles instead of 18.

Instr. No	Instructions	Cycles									
		1	2	3	4	5	6	7	8	9	10
0	MUL R5, R6, R7	F	D	E1	E2	E3	E4	M	W		
3	MUL R4, R7, R7		F	D	E1	E2	E3	E4	M	W	
2	ADD R5, R5, R6			F	D	-	-	E1	E2	E3	M
4	ADD R6, R7, R5				F	-	-	D	-	-	E1
6	ADD R7, R1, R4							F	-	-	D
5	ADD R3, R0, R6										F
		11	12	13	14	15	16	17	18	19	20
0	MUL R5, R6, R7										
3	MUL R4, R7, R7										
2	ADD R5, R5, R6	W									
4	ADD R6, R7, R5	E2	E3	M	W						
6	ADD R7, R1, R4	E1	E2	E3	M	W					
5	ADD R3, R0, R6	D	-	E1	E2	E3	M	W			

5 Tomasulo's Algorithm

Remember that Tomasulo's algorithm requires tag broadcast and comparison to enable wake-up of dependent instructions. In this question, we will calculate the number of tag comparators and size of tag storage required to implement Tomasulo's algorithm in a machine that has the following properties:

- 8 functional units where each functional unit has a dedicated separate tag and data broadcast bus
- 32 64-bit architectural registers
- 16 reservation station entries per functional unit
- Each reservation station entry can have two source registers

Answer the following questions. Show your work for credit.

- (a) What is the number of tag comparators per reservation station entry?

$$8 * 2$$

- (b) What is the total number of tag comparators in the entire machine?

$$16 * 8 * 2 * 8 + 8 * 32$$

- (c) What is the (minimum possible) size of the tag?

$$\log(16 * 8) = 7$$

- (d) What is the (minimum possible) size of the register alias table (or, frontend register file) in bits?

$$72 * 32 \text{ (64 bits for data, 7 bits for the tag, 1 valid bit)}$$

- (e) What is the total (minimum possible) size of the tag storage in the entire machine in bits?

$$7 * 32 + 7 * 16 * 8 * 2$$

6 Tomasulo's Algorithm - Reverse Engineering

In this problem, we will give you the state of the Register Alias Table (RAT) and Reservation Stations (RS) for an out-of-order execution engine that employs Tomasulo's algorithm, as we discussed in lectures. Your job is to determine the original sequence of **four instructions** in program order.

The out-of-order machine in this problem behaves as follows:

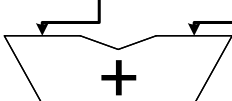
- The frontend of the machine has a one-cycle fetch stage and a one-cycle decode stage. The machine can fetch one instruction per cycle, and can decode one instruction per cycle.
- The machine executes *only* register-type instructions, e.g., $OP\ R_{dest} \leftarrow R_{src1},\ R_{src2}$.
- The machine dispatches one instruction per cycle into the reservation stations, in program order. Dispatch occurs during the decode stage.
- An instruction always allocates the first reservation station that is available (in top-to-bottom order) at the required functional unit.
- When an instruction in a reservation station finishes executing, the reservation station is cleared.
- The adder and multiplier **are not** pipelined. An add operation takes 2 cycles. A multiply operation takes 3 cycles.
- The result of an addition and multiplication is broadcast to the reservation station entries and the RAT in the writeback stage. A dependent instruction can begin execution in the next cycle after the writeback if it has all of its operands available in the reservation station entry.
- When multiple instructions are ready to execute at a functional unit at the same cycle, the oldest ready instruction is chosen to be executed first.

Initially, the machine is empty. Four instructions then are fetched, decoded, and dispatched into reservation stations. Pictured below is the state of the machine when the final instruction has been dispatched into a reservation station:


RAT

Reg	V	Tag	Value
R0	—	—	—
R1	0	A	5
R2	1	—	8
R3	0	E	—
R4	0	B	—
R5	—	—	—

ID	V	Tag	Value	V	Tag	Value
A	0	D	—	1	—	8
B	0	A	—	0	A	—
C	—	—	—	—	—	—



ID	V	Tag	Value	V	Tag	Value
D	1	—	5	1	—	5
E	0	A	—	0	B	—
F	—	—	—	—	—	—



- (a) Give the four instructions that have been dispatched into the machine, in program order. The source registers for the first instruction can be specified in either order. Give instructions in the following format: “opcode destination \leftarrow source1, source2.”

MUL	R1	\leftarrow	R1	,	R1
ADD	R1	\leftarrow	R1	,	R2
ADD	R4	\leftarrow	R1	,	R1
MUL	R3	\leftarrow	R1	,	R4

- (b) Now assume that the machine flushes all instructions out of the pipeline and restarts fetch from the first instruction in the sequence above. Show the full pipeline timing diagram below for the sequence of four instructions that you determined above, from the fetch of the first instruction to the writeback of the last instruction. Assume that the machine stops fetching instructions after the fourth instruction.

As we saw in lectures, use “F” for fetch, “D” for decode, “En” to signify the nth cycle of execution for an instruction, and “W” to signify writeback. You may or may not need all columns shown.

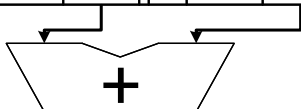
Cycle:	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16
MUL R1 \leftarrow R1, R1	F	D	E1	E2	E3	W										
ADD R1 \leftarrow R1, R2		F	D				E1	E2	W							
ADD R4 \leftarrow R1, R1			F	D						E1	E2	W				
MUL R3 \leftarrow R1, R4				F	D								E1	E2	E3	W

- (c) Finally, show the state of the RAT and reservation stations at the end of the **12th cycle** of execution in the figure below. Complete all blank parts.

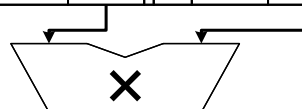
RAT

Reg	V	Tag	Value
R0	—	—	—
R1	1	—	33
R2	1	—	8
R3	0	E	—
R4	1	—	66
R5	—	—	—

ID	V	Tag	Value	V	Tag	Value
A	—	—	—	—	—	—
B	—	—	—	—	—	—
C	—	—	—	—	—	—



ID	V	Tag	Value	V	Tag	Value
D	—	—	—	—	—	—
E	1	—	33	1	—	66
F	—	—	—	—	—	—



7 Out-of-Order Execution

In this problem, we consider an in-order fetch, out-of-order dispatch, and in-order retirement execution engine that employs Tomasulo's algorithm. This engine behaves as follows:

- The engine has four main pipeline stages: Fetch (F), Decode (D), Execute (E), and Write-back (W).
- The engine can fetch one instruction per cycle, decode one instruction per cycle, and write back the result of one instruction per cycle.
- The engine has two execution units: 1) an adder for executing ADD instructions and 2) a multiplier for executing MUL instructions.
- The execution units are fully pipelined. The adder has two stages (E1-E2) and the multiplier has four stages (E1-E2-E3-E4). Execution of each stage takes one cycle.
- The adder has a two-entry reservation station and the multiplier has a four-entry reservation station.
- An instruction always allocates the first available entry of the reservation station (in top-to-bottom order) of the corresponding execution unit.
- Full data forwarding is available, i.e., during the last cycle of the E stage, the tags and data are broadcast to the reservation station and the Register Alias Table (RAT). For example, an ADD instruction updates the reservation station entries of the dependent instructions in E2 stage. So, the updated value can be read from the reservation station entry in the next cycle. Therefore, a dependent instruction can potentially begin its execution in the next cycle (after E2).
- The multiplier and adder have separate output data buses, which allow both the adder and the multiplier to update the reservation station and the RAT in the same cycle.
- An instruction continues to occupy a reservation station slot until it finishes the Write-back (W) stage. The reservation station entry is deallocated after the Write-back (W) stage.

7.1 Problem Definition

The processor is about to fetch and execute *six* instructions. Assume the *reservation stations (RS)* are all initially empty and the initial state of the *register alias table (RAT)* is given below in Figure (a). Instructions are fetched, decoded and executed as discussed in class. At some point during the execution of the six instructions, a snapshot of the state of the RS and the RAT is taken. Figures (b) and (c) show the state of the RS and the RAT at the snapshot time. A dash (–) indicates that a value has been cleared. A question mark (?) indicates that a value is unknown.

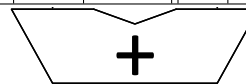
Reg	Valid	Tag	Value
R0	1	–	1900
R1	1	–	82
R2	1	–	1
R3	1	–	3
R4	1	–	10
R5	1	–	5
R6	1	–	23
R7	1	–	35
R8	1	–	61
R9	1	–	4

Initial state of the RAT

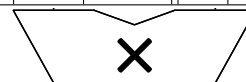
Reg	Valid	Tag	Value
R0	1	?	1900
R1	0	Z	?
R2	1	?	12
R3	1	?	3
R4	1	?	10
R5	0	B	?
R6	1	?	23
R7	0	H	?
R8	1	?	350
R9	0	A	?

Snapshot state of the RAT

ID	V	Tag	Value	V	Tag	Value
A	1	?	350	1	?	12
B	0	A	?	0	Z	?

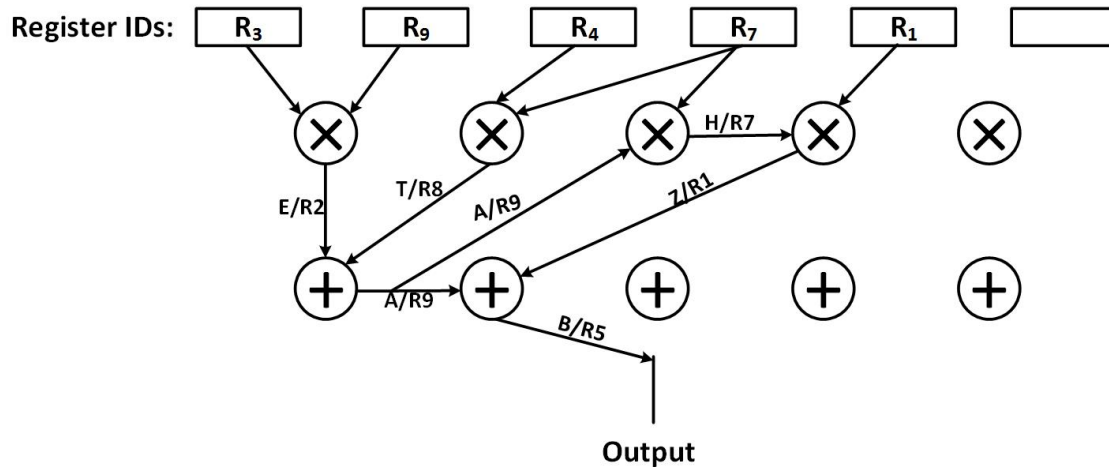


ID	V	Tag	Value	V	Tag	Value
–	–	–	–	–	–	–
T	1	?	10	1	?	35
H	1	?	35	0	A	?
Z	1	?	82	0	H	?



7.2 (a) Data Flow Graph

Based on the information provided above, identify the instructions and complete the dataflow graph below for the six instructions that have been fetched. Please appropriately connect the nodes using edges and specify the direction of each edge. Label each edge with the destination architectural register and the corresponding Tag. *Note that you may **not** need to use all registers and/or nodes provided below.*



7.3 (b) Program Instructions

Fill in the blanks below with the six-instruction sequence in program order. When referring to registers, please use their architectural names (R0 through R9). Place the register with the smaller architectural name on the left source register box. For example, `ADD R8 \leftarrow R1, R5`.

MUL	R2	\leftarrow	R3	,	R9
MUL	R8	\leftarrow	R4	,	R7
ADD	R9	\leftarrow	R2	,	R8
MUL	R7	\leftarrow	R7	,	R9
MUL	R1	\leftarrow	R1	,	R7
ADD	R5	\leftarrow	R1	,	R9

8 Out-of-Order Execution - Reverse Engineering

A five instruction sequence executes according to Tomasulo's algorithm. Each instruction is of the form ADD DR,SR1,SR2 or MUL DR,SR1,SR2. ADDs are pipelined and take 9 cycles (F-D-E1-E2-E3-E4-E5-E6-WB). MULs are also pipelined and take 11 cycles (two extra execute stages). An instruction must wait until a result is in a register before it sources it (reads it as a source operand). For instance, if instruction 2 has a read-after-write dependence on instruction 1, instruction 2 can start executing in the next cycle after instruction 1 writes back (shown below).

```
instruction 1    |F|D|E1|E2|E3|.....|WB|
instruction 2    |F|D|-|-|.....|-|E1|
```

The machine can fetch one instruction per cycle, and can decode one instruction per cycle.

The register file before and after the sequence are shown below.

	Valid	Tag	Value
R0	1		4
R1	1		5
R2	1		6
R3	1		7
R4	1		8
R5	1		9
R6	1		10
R7	1		11

	Valid	Tag	Value
R0	1		310
R1	1		5
R2	1		410
R3	1		31
R4	1		8
R5	1		9
R6	1		10
R7	1		21

- (a) Complete the five instruction sequence in program order in the space below. Note that we have helped you by giving you the opcode and two source operand addresses for the fourth instruction. (The program sequence is unique.)

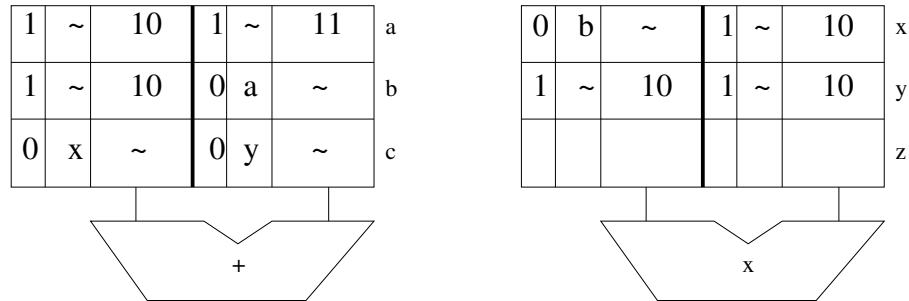
Give instructions in the following format: "opcode destination \Leftarrow source1, source2."

ADD	R7	\Leftarrow	R6	,	R7
ADD	R3	\Leftarrow	R6	,	R7
MUL	R0	\Leftarrow	R3	,	R6
MUL	R2	\Leftarrow	R6	,	R6
ADD	R2	\Leftarrow	R0	,	R2

- (b) In each cycle, a single instruction is fetched and a single instruction is decoded.

Assume the reservation stations are all initially empty. Put each instruction into the next available reservation station. For example, the first ADD goes into “a”. The first MUL goes into “x”. Instructions remain in the reservation stations until they are completed. Show the state of the reservation stations at the end of cycle 8.

Note: to make it easier for the grader, when allocating source registers to reservation stations, please always have the higher numbered register be assigned to source2.



- (c) Show the state of the Register Alias Table (Valid, Tag, Value) at the end of cycle 8.

	Valid	Tag	Value
R0	0	x	4
R1	1	~	5
R2	0	c	6
R3	0	b	7
R4	1	~	8
R5	1	~	9
R6	1	~	10
R7	0	a	11